


A 3D SRAM tier on a B200-class GPU: C2, 2 layers, 4.22 GB (geometric upper bound: full 800 mm² placement per layer per die) on Llama-3.1-8B decode, B=8, N=2048


All three rates are bytes a kernel consumes per unit time, not link pin rates. HBM: logical bytes / measured B200 decode step time, fitted slope over four anchors.

L2: Vulkan compute-shader benchmark, 21 TB/s within a partition and 16.8 across. 3D: device-team deck gives 17.7 read-oriented (3.37x derate calibrated on H100 L2); the read/fill mix table is from a Slack reply, not the deck.

A Where the tier sits



B What each rung of modeling refinement does



Model calculation, not a measurement. BW(r) = 15 / 9.4 / 3.8 at r = 1.0 / 0.8 / 0.6 comes from a device-team Slack reply (2026-09-20), not the audited deck; its derivation is unconfirmed.

The cache rungs therefore hold only if that table holds: the serial bound (solid) needs no such table, the overlapped bound (faint) rests on BW(0.6) = 3.8 < 6.4. B200 L2 is not a counterexample: its 19 TB/s exceeds HBM.