


Read delivery at this point: array peak 904 TB/s (40,353 macros x 22.4 GB/s) far exceeds the fabric, so $B_R = \min(\text{array}/3.37, \text{L2-to-SM}) = 19 \text{ TB/s}$. The deck's 17.7 was the 320 mm² array. Everything on the x axis is an architectural assumption we sweep. Solid = overlapped bound, faint = serial bound. Shared-port service $T = R/B_R + W/(\beta B_R)$. HBM effective 6.4 TB/s

How much write capability makes a demand-filled 3D SRAM cache viable? C2 2-layer 4.22 GB (800 mm² full-placement upper bound), Llama-



Model calculation, not a measurement. Tag-first: hit reads, miss fills ($R = hD$, $W = (1-h)D$). Parallel: every access reads the data array ($R = D$).

The serial bound never breaks even for $\beta \leq 1$. Slack-implied beta re-expresses the 2026-09-20 reply ($r = 0.8 \rightarrow 9.4$, $r = 0.6 \rightarrow 3.8 \text{ TB/s}$) in this model; it is not a deck value.